

assignment4

April 27, 2020

1 Machine Learning and Computer Vision

1.1 Assignment 4

This assignment contains 2 programming exercises.

1.2 Problem 1: Canny Edge Detection

In this problem, you are required to write a function that performs Canny Edge Detection¹. The function has the following specifications:

- It takes in two inputs: a grayscale image (geisel.jpg), and a threshold t_e .

- It returns the edge image.

- You are allowed the use of loops.

A brief description of the algorithm is given below. Make sure your function reproduces each step as given.

(i) Smoothing:

It is inevitable that all images taken from a camera will contain some amount of noise. To prevent noise from being mistaken for edges, noise must be reduced. Therefore the image is first smoothed by applying a Gaussian filter. A Gaussian kernel with standard deviation $\sigma = 1.4$ (shown below) is to be used.

You can perform this filtering step by using the `scipy.signal.convolve2d()` function.

(ii) Finding Gradients:

The next step is to find the horizontal and vertical gradients of the smoothed image using the Sobel operators. The gradient images in the x and y -direction, G_x and G_y are found by applying the kernels k_x and k_y given below. These operations can be performed using `scipy.signal.convolve2d()` in the same manner as before.

The corresponding gradient magnitude image is computed using:

and the edge direction image is calculated as follows:

(iii) Non-maximum Suppression (NMS):

The purpose of this step is to convert the thick edges in the gradient magnitude image to “sharp” edges. This is done by preserving all local maxima in the gradient image, and deleting everything else. This is carried out by recursively performing the following steps for each pixel in the gradient image:

Round the gradient direction to nearest 45, corresponding to the use of an 8-connected neighbourhood.

Compare the edge strength of the current pixel with the edge strength of the pixel in the positive and negative gradient direction i.e. if the gradient direction is north (= 90), then compare with the pixels to the north and south.

If the edge strength of the current pixel is largest; preserve the value of the edge strength. If not, suppress (remove) the value.

(iv) Thresholding:

The edge-pixels remaining after the NMS step are (still) marked with their strength. Many of these will probably be true edges in the image, but some may be caused by noise or color variations. The simplest way to remove these would be to use a threshold, so that only edges stronger than a certain value would be preserved. Use the input *te* to perform thresholding on the non-maximum suppressed magnitude image.

Evaluate your canny edge detection function on *geisel.jpg* for a suitable value of *te* that retains the structural edges, and removes the noisy ones.

Things to turn in:

Image after smoothing, the original gradient magnitude image, the image after NMS, and the final edge image after thresholding.

The value for *te* that you used to produce the final edge image.

```
[54]: import numpy as np
from imageio import imread
import matplotlib.pyplot as plt
from scipy import signal
def rgb2gray(rgb):
    grayimg = rgb[:, :, 0] * 0.299 + rgb[:, :, 1] * 0.587 + rgb[:, :, 2] * 0.114
    return grayimg

#Canny Edge Detection Function
def canny_edge(img, te):
    k = 1/159 * np.array([[2,4,5,4,2],
                          [4,9,12,9,4],
                          [5,12,15,12,5],
                          [4,9,12,9,4],
                          [2,4,5,4,2]])

    smothing = signal.convolve2d(img, k, boundary='symm', mode='same')

    kx = np.array([[-1, 0, 1],
                  [-2, 0, 2],
                  [-1, 0, 1]])

    ky = kx.T
    gx = signal.convolve2d(smothing, kx, boundary='symm', mode='same')
    gy = signal.convolve2d(smothing, ky, boundary='symm', mode='same')
```

```

g = np.sqrt(gx*gx + gy*gy)
gdir = np.arctan(gy/gx)
gdir = np.round(gdir*4/np.pi)+3
gg = np.copy(g)

for i in range(g.shape[0]):
    for j in range(g.shape[1]):
        p, q = i, j
#         north
        while p >= 0 and p < g.shape[0] and q >= 0 and q < g.shape[1] and
→(gdir[p, q] == 5 or gdir[p, q] == 1):
            if g[p, q] > gg[i, j]:
                gg[i, j] = 0
            p, q = p+1, q
#         south
        while p >= 0 and p < g.shape[0] and q >= 0 and q < g.shape[1] and
→(gdir[p, q] == 1 or gdir[p, q] == 5):
            if g[p, q] > gg[i, j]:
                gg[i, j] = 0
            p, q = p-1, q
#         east
        while p >= 0 and p < g.shape[0] and q >= 0 and q < g.shape[1] and
→gdir[p, q] == 3:
            if g[p, q] > gg[i, j]:
                gg[i, j] = 0
            p, q = p, q+1
#         west
        while p >= 0 and p < g.shape[0] and q >= 0 and q < g.shape[1] and
→gdir[p, q] == 3:
            if g[p, q] > gg[i, j]:
                gg[i, j] = 0
            p, q = p, q-1
#         northeast
        while p >= 0 and p < g.shape[0] and q >= 0 and q < g.shape[1] and
→gdir[p, q] == 4:
            if g[p, q] > gg[i, j]:
                gg[i, j] = 0
            p, q = p+1, q+1
#         southwest
        while p >= 0 and p < g.shape[0] and q >= 0 and q < g.shape[1] and
→gdir[p, q] == 4:
            if g[p, q] > gg[i, j]:
                gg[i, j] = 0
            p, q = p-1, q-1
#         southeast

```

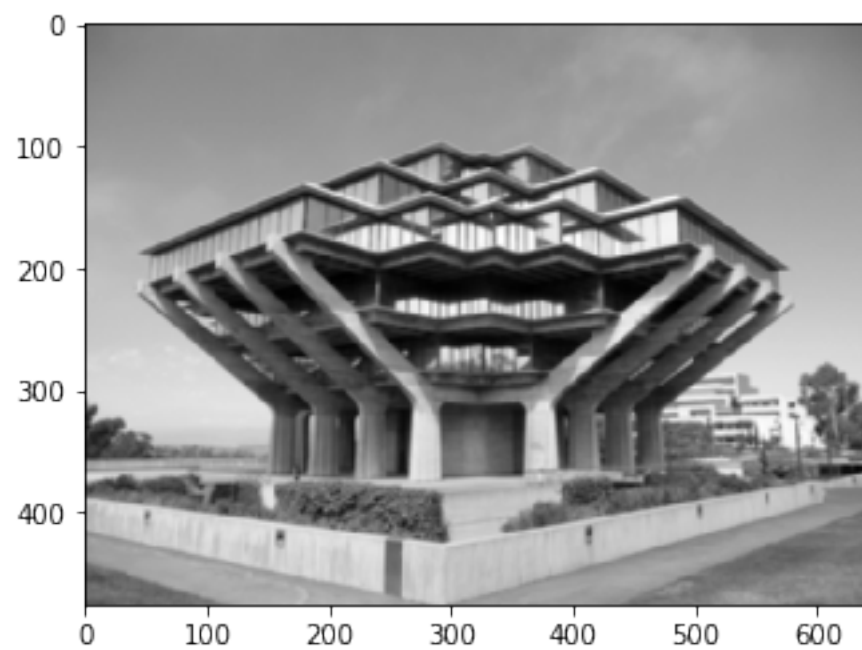
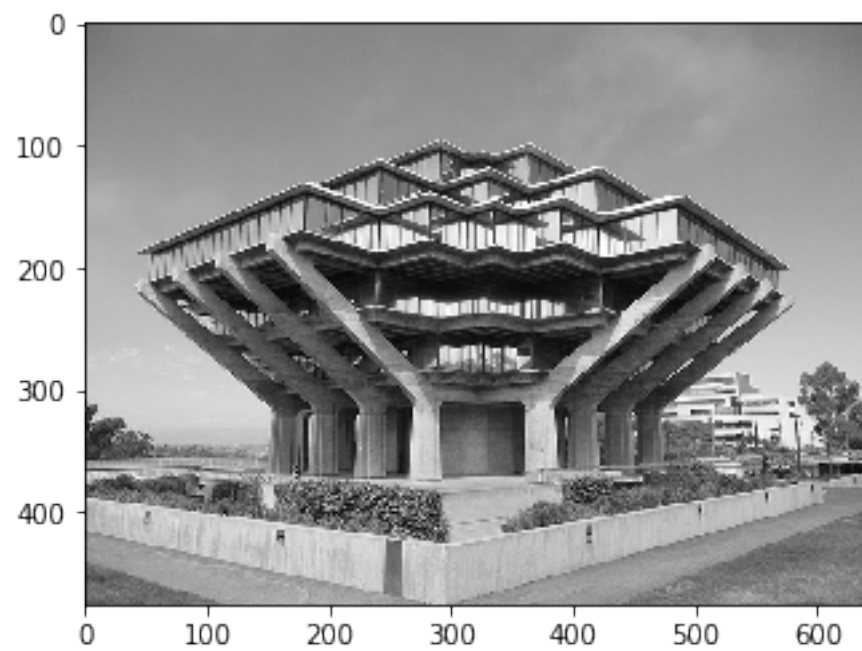
```

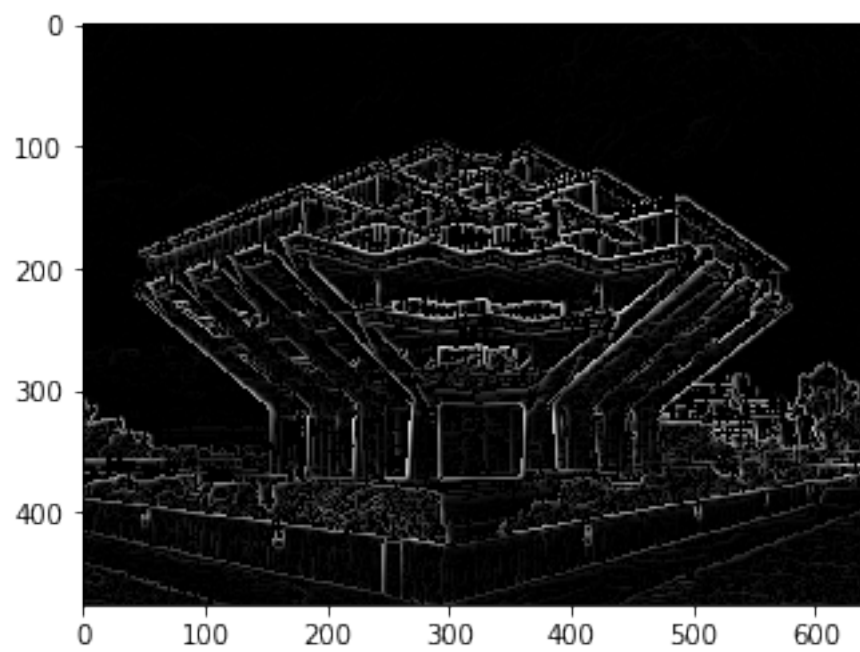
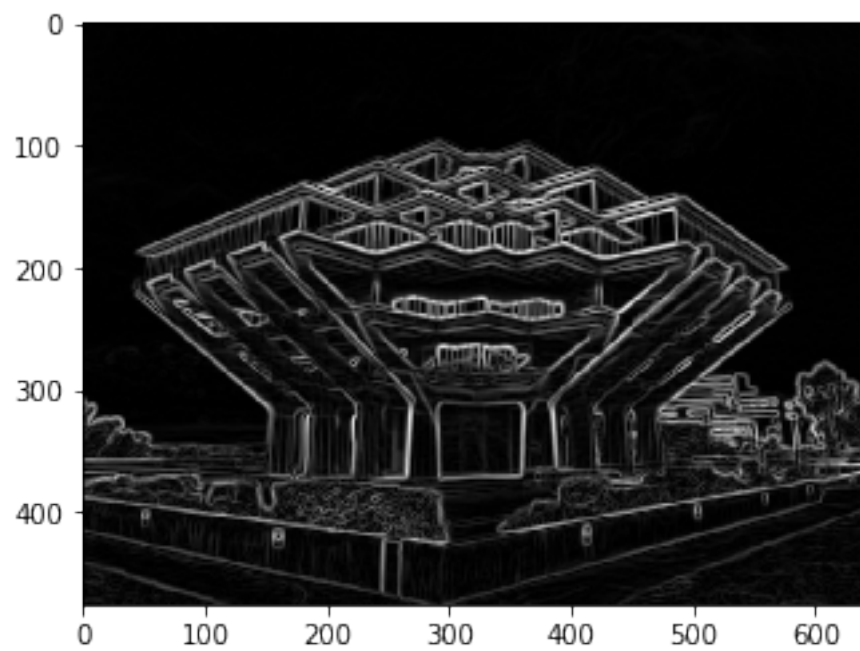
        while p >= 0 and p < g.shape[0] and q >=0 and q < g.shape[1] and
→gdir[p, q] == 2:
            if g[p, q]>gg[i, j]:
                gg[i,j]=0
                p, q = p-1, q+1
#           northwest
        while p >= 0 and p < g.shape[0] and q >=0 and q < g.shape[1] and
→gdir[p, q] == 2:
            if g[p, q]>gg[i, j]:
                gg[i,j]=0
                p, q = p-1, q+1
            detected_img = np.copy(gg)
            for i in range(img.shape[0]):
                for j in range(img.shape[1]):
                    if gg[i, j] > te:
                        detected_img[i, j] = 1
                    else:
                        detected_img[i, j] = 0
            return smothing, g, gg, detected_img
#         return gdir
#         return detected_img

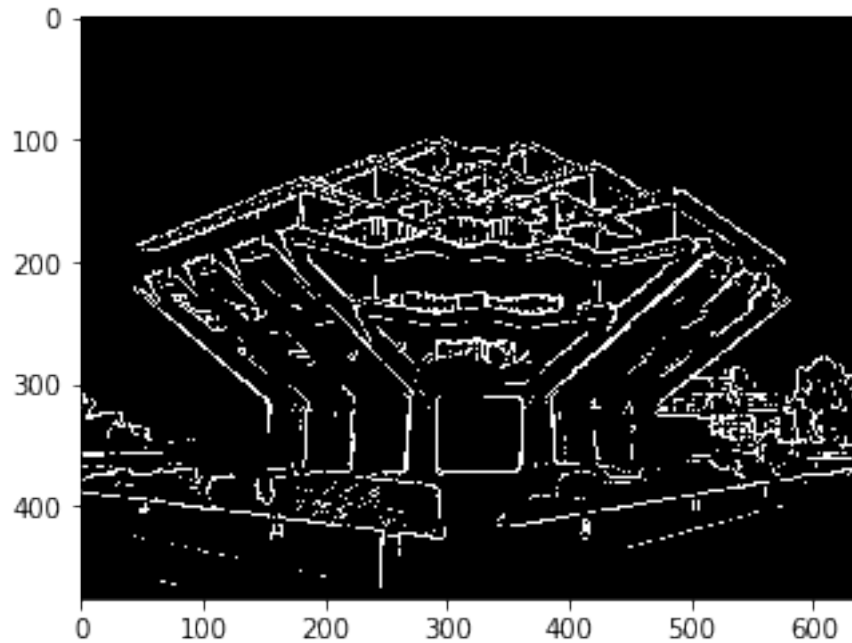
#Import image here
im = imread('geisel.jpg')
te = 120
grayim = rgb2gray(im)
plt.imshow(grayim, cmap='gray')

smothingim, g, gg, detected_img = canny_edge(grayim, te)
plt.figure()
plt.imshow(smothingim, cmap='gray')
plt.figure()
plt.imshow(g, cmap='gray')
plt.figure()
plt.imshow(gg, cmap='gray')
plt.figure()
plt.imshow(detected_img, cmap='gray')
plt.show()
print('threshold = %d' %te)

```







threshold = 120

1.3 Problem 2 Adaptive Histogram Equalization

It is often found in image processing and related fields that real world data is unsuitable for direct use. This warrants the inclusion of pre-processing steps before any other operations are performed. An example of this is histogram equalization (HE) and its extension adaptive histogram equalization (AHE). The goal of this problem is to implement a function for AHE as described in Chapter 1 of Adaptive Histogram Equalization - A Parallel Implementation². The function has the following specifications:

- (i) The desired function `AHE()` takes two inputs: the image ("beach.png") `im` and the contextual region size `win_size`.
- (ii) Using the pseudocode in Algorithm as a reference, compute the enhanced image after AHE.
- (iii) You may use loops if necessary. You should not make use of any inbuilt functions for AHE or HE.
- (iv) The function returns one output: the enhanced image after AHE.

Evaluate your function on the image `beach.png` for `win_size` = 33, 65 and 129. In your report, include the original image, the 3 images after AHE. Make sure to resize all images to ensure they do not take up too much space. Additionally, include your answers (no more than three sentences each) to the following questions:

How does the original image qualitatively compare to the images after AHE?

```

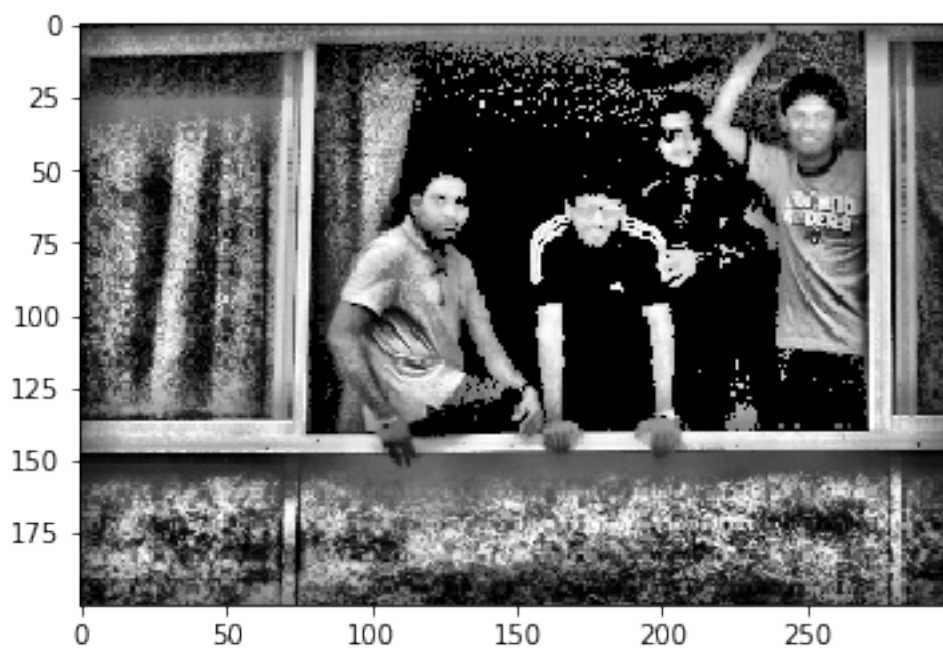
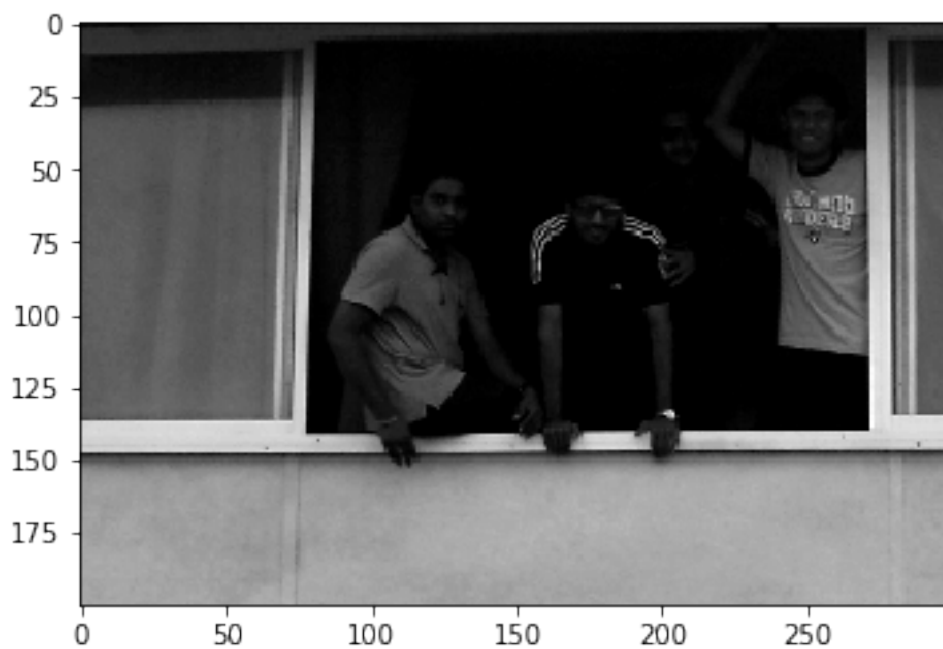
[73]: from imageio import imread
import numpy as np
import matplotlib.pyplot as plt

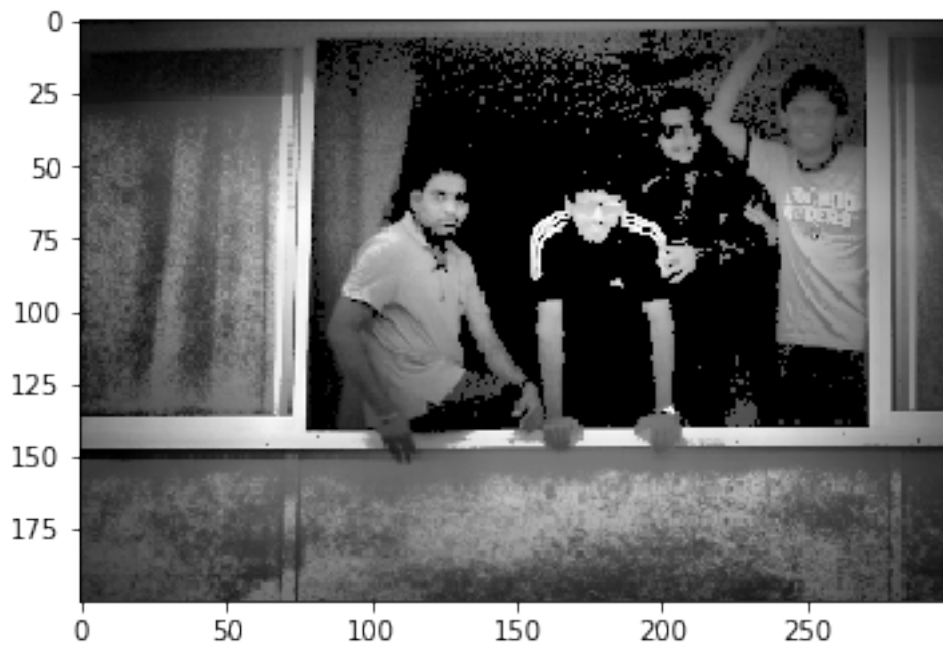
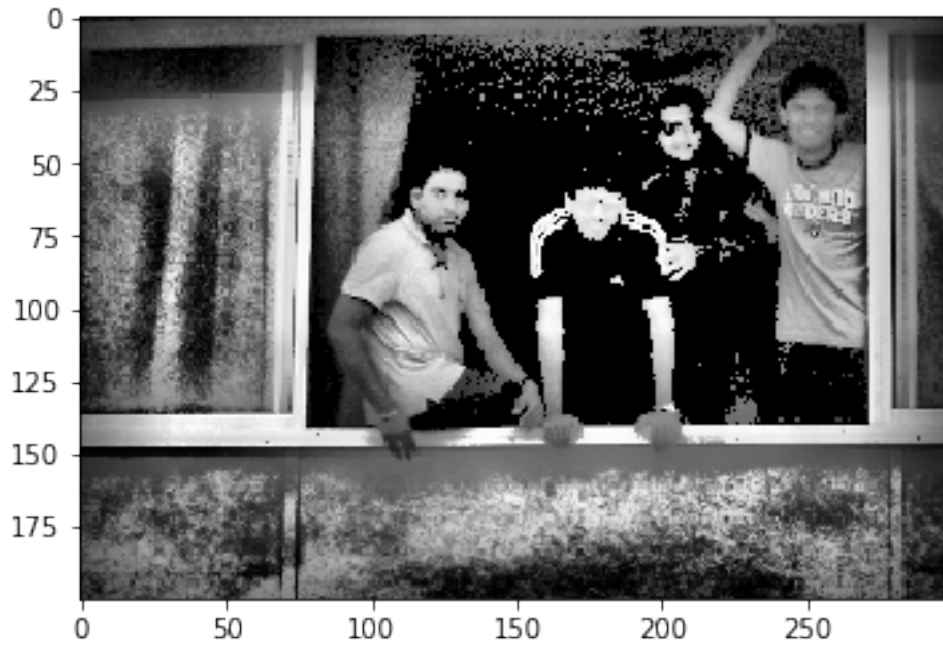
def AHE(im, win_size):
    contextual_regin = np.zeros((win_size, win_size))
    output = np.zeros(im.shape)
    for x in range(im.shape[0]):
        for y in range(im.shape[1]):
            rank = 0
            for i in range(-(win_size-1)//2, (win_size+1)//2):
                for j in range(-(win_size-1)//2, (win_size+1)//2):
                    if 0<=x+i and x+i < im.shape[0] and 0 <= y+j and y+j < im.
→shape[1] and im[x, y] > im[x+i, y+j]:
                        rank += 1
            output[x, y] = rank
    output = output*255/(win_size*win_size)
    return output

#Import image here
#Sample call
#Plotting code below
im = imread('beach.png')
plt.imshow(im[200:400, 200:500], cmap='gray')

plt.figure()
output1 = AHE(im[200:400, 200:500], 33)
plt.imshow(output1, cmap='gray')
plt.figure()
output2 = AHE(im[200:400, 200:500], 65)
plt.imshow(output2, cmap='gray')
plt.figure()
output3 = AHE(im[200:400, 200:500], 129)
plt.imshow(output3, cmap='gray')
plt.show()

```



1.4 Conclusion

Have you accomplished all parts of your assignment? What concepts did you use or learn in this assignment? What difficulties have you encountered? Explain your result for each section.

Please write one or two short paragraphs in the below Markdown window (double click to edit).

*** Your Conclusion: ***

– Problem 1: According to the requirements, four images are shown: image after smoothing, the original gradient magnitude image, the image after NMS, and the final edge image after thresholding and the thresholding is 120. The function is completed by following instructions step-by-step. Since the instructions are detailed, I didn't encounter any difficulties.

Problem 2: the pseudocode is given, so the function is not difficult to create. It takes a lot of time for the functions to give a result image, because for each pixel in a image, it needs to compare $\text{win_size} * \text{win_size}$ times to get the rank. For $\text{window_size}=169$, it needs to compare 16,641 times to get one rank... It is super slow. So, I sliced out the middle of the image where there are four people. According to my results, images with smaller window size are more qualitative. The faces of the people in the middle of the image are not dark in the shadow and the curtains are more clear when the window size is small.

** Submission Instructions **

Remember to submit your pdf version of this notebook to Gradescope. You can find the export option at File → Download as → PDF via LaTeX